

COLL 150

Artificial Intelligence in Speculative Media and Literatures

INSTRUCTOR	Ashley Y. Gao, Department of Computer Science
EMAIL	ygao18@wm.edu
OFFICE	Integrated Science Center 4, Office 2381
OFFICE HOURS	[TBD], and by appointment
MEETING TIME	[TBD]
CLASSROOM	[TBD]
CREDITS	3
PREREQUISITES	None. No programming, mathematics, or literature background is assumed or required.

A note before you read further. This is a Computer Science course that assigns Ovid, Marlowe, and Mary Shelley, and it is a literature course that will ask you to read a loss function.

If you are a prospective CS major who has never written a close reading, and if you are a prospective English major who has never seen a partial derivative, you are both exactly where you belong. Everything technical will be taught from zero.

COURSE DESCRIPTION

Long before anyone wrote a line of code, people were telling stories about made minds: an ivory statue that warms under a sculptor's hands, a creature assembled from corpses and abandoned by its maker, a teapot that opens its eyes after a hundred years of use. Today we build systems that we describe, without quite meaning it, in the same vocabulary — systems that “learn,” “hallucinate,” “want,” “collapse,” “forget.”

This seminar takes that coincidence seriously. Over thirteen weeks we will read speculative literature, film, and games alongside the actual technical machinery of modern machine learning, and we will treat both as *texts* — as artifacts made by people, encoding assumptions, arguments, and desires. A loss function is a mathematical object; it is also a sentence about what is worth minimizing. Mode collapse is a failure mode of a generative model; it is also a fable about what happens to a system rewarded for producing what pleases. We will move back and forth between the two registers until the border between them starts to feel less obvious than it did in August.

The organizing question of the course is not “can machines think?” It is a set of harder, stranger questions. **What do we want from a made mind, and what does the shape of that wanting tell us about ourselves?** Who counts as a maker, who counts as a creature, and who gets to decide? What does it cost to be the thing that answers?

WHAT THIS COURSE IS (COLL 150)

COLL 150 is a first-year seminar in W&M's College Curriculum. It is a **writing course** with a topic, not a topics course with some writing attached. Roughly a third of our time together will be spent explicitly on the craft of scholarly work: framing a question, finding and evaluating sources, structuring an argument, giving and receiving feedback, and revising — really revising, not just proofreading.

By the end of the semester, you will be able to:

1. **Analyze and interpret** speculative texts and technical artifacts closely, attending to form, structure, rhetoric, and the assumptions each embeds. (*Critical Thinking*)
2. **Work independently** to understand unfamiliar material — a nineteenth-century novel, a 2014 machine learning paper — and form and defend your own judgments about it. (*Independent Inquiry*)
3. **Write scholarly prose** that communicates a complex, contestable argument in your own voice, supported by evidence and properly cited. (*Scholarly Writing*)
4. **Locate, evaluate, and cite** scholarly and technical sources, including through Swem Library's research tools. (*Information Literacy*)
5. **Speak and listen** as a member of a seminar: pose real questions, disagree productively, and present research to an audience. (*Oral Communication*)
6. **Explain the core mechanics** of contemporary machine learning — objective functions, gradient descent, generative modeling, reinforcement learning — accurately, in plain English, to a non-specialist.

A grade of **C– or better** fulfills the COLL 150 requirement.

REQUIRED MATERIALS

Please acquire the editions listed below where an edition is specified; page numbers in class will refer to them. Everything not listed here will be posted to **Blackboard** as a PDF or a licensed streaming link. Cost is a real constraint — if any of this is a hardship, **email me in the first week** and we will solve it quietly and without fuss. Swem holds copies of everything, and I keep a small lending shelf.

Books to buy

- Ovid, *Metamorphoses*, trans. Stanley Lombardo (Hackett) — or any complete verse translation you prefer.
- Mary Shelley, *Frankenstein* — **1818 text**, Norton Critical Edition, 2nd ed. (We will discuss why the edition matters.)
- Christopher Marlowe, *Doctor Faustus* — A-text and B-text; Norton Critical or Revels Student Edition.
- Steven Connor, *The Madness of Knowledge: On Wisdom, Ignorance and Fantasies of Knowing*.
- Ted Chiang, *Exhalation* (Vintage) — for “The Lifecycle of Software Objects.”

- Kazuo Ishiguro, *Klara and the Sun* (Vintage).
- Toriyama Sekien, *Japandemonium Illustrated: The Yokai Encyclopedias of Toriyama Sekien*, trans. Hiroko Yoda & Matt Alt (Dover).
- Karel Čapek, *R.U.R.* — any translation; public-domain versions are free online.

Films

The Matrix (Wachowskis, 1999) · *Ex Machina* (Garland, 2014) · *Her* (Jonze, 2013). Optional: *Ghost in the Shell* (Oshii, 1995), *Blade Runner* (Scott, 1982).

Games

Detroit: Become Human (Quantic Dream, 2018) and the *Mass Effect* trilogy (BioWare, 2007–2012).

You are not required to buy or complete these games. Games are long and expensive, and I will not pretend otherwise. Playing is a pleasure, not a prerequisite. Interactive sequences that we need to discuss will be provided as YouTube links available on the course website.

Technical readings

All free and online; several are videos. These are chosen for accessibility, not rigor-for-its-own-sake. Full links on Blackboard.

- 3Blue1Brown, *Neural Networks*, chs. 1–4 (video)
- Sutton & Barto, *Reinforcement Learning: An Introduction*, ch. 1 (free PDF)
- Goodfellow et al., “Generative Adversarial Networks” (2014) — abstract, Figure 1, and §3 only
- Amodei et al., “Concrete Problems in AI Safety” (2016) — §§1–3
- Krakovna et al., “Specification Gaming: The Flip Side of AI Ingenuity,” with the accompanying public list of examples
- Turing, “Computing Machinery and Intelligence” (*Mind*, 1950)
- Bender, Gebru, McMillan-Major & Shmitchell, “On the Dangers of Stochastic Parrots” (2021)
- Ted Chiang, “ChatGPT Is a Blurry JPEG of the Web” (*The New Yorker*, 2023)

Content advisory. Our texts include depictions of violence, bodily horror, suicide, sexual coercion, abuse (including the abuse of children and of android characters coded as children), racialized dehumanization, and self-destruction. *Frankenstein*, *Doctor Faustus*, and *Detroit: Become Human* are each, in their way, difficult. I will flag specific material in advance, and there are always alternate paths through an assignment. Come talk to me.

ASSIGNMENTS AND GRADING

ASSIGNMENT	WEIGHT	DUE
Essay 1: Close Reading (1,200–1,500 words) + required revision	25%	Draft Unit 5 · Revision Unit 7

ASSIGNMENT	WEIGHT	DUE
Essay 2: Reading the Machine (1,500–1,800 words)	25%	Draft Unit 9 · Final Unit 11
Midterm: Creative Writing Response	10%	Exam Period
Research project: final essay (2,500–3,000 words)	20%	Exam period
Symposium presentation (8 min + Q&A)	10%	Nov 30 – Dec 4
Final Exam: Creative Writing Response	10%	Exam Period

Total polished prose: roughly 2,000–4,000 words. That is the COLL 150 workload, and it is front-loaded with support: no major assignment goes from blank page to grade without a draft, a workshop, and a conference with me.

Essay 1 — Close Reading

Take a passage of no more than thirty lines (or one page, or one scene, or one game sequence under three minutes) from Units 1–4 and make an argument about how it *works*. Not what it means in the abstract; how the thing is built such that it means that. This is the foundational move of literary analysis and, done honestly, it is much harder than it sounds. Full revision required; **I grade the revision.**

Essay 2 — Reading the Machine

Choose a technical artifact — a loss function, an evaluation benchmark, a reward specification, a model card, a dataset's documentation, a paper's Figure 1 — and read it the way you read Ovid in Essay 1. What argument does it make? What does it take for granted? Who is imagined as its user, and who is imagined as its object? What does it exclude in order to be well-formed?

Research Project

A semester-long argument of your own devising, on any question this course has provoked and that you can defend with evidence. We build it in stages: a question (Unit 9), a proposal and annotated bibliography (Unit 10, following our Swem Library session), a full draft workshopped in the week of November 23, a conference with me, and a final essay. You will present it at the **Symposium** in the final week.

Grading scale

A 93–100 · A– 90–92 · B+ 87–89 · B 83–86 · B– 80–82 · C+ 77–79 · C 73–76 · C– 70–72 · D+ 67–69 · D 63–66 · D– 60–62 · F below 60

Late work

Every student has **three no-questions-asked late days** to spend across the semester on major assignments — not on reading responses, which are worthless if they arrive after the discussion they were meant to prepare. Use them however you like; just tell me you are using one. Beyond that, work drops one-third of a letter grade per day. If something serious is happening in your life, do not spend your late days silently — come talk to me and we will make a plan. I would much rather hear from you early than read an apology later.

Attendance

This is a seminar. Its entire method is that you are in the room with each other. Two unexcused absences carry no penalty; each subsequent one lowers your participation grade. Miss more than five sessions and we will need to talk about whether you can pass. Illness, religious observance, athletic and academic travel, and family emergencies are excused — just notify me when you can.

POLICY ON GENERATIVE AI

We are going to spend a semester studying these systems. It would be strange to pretend they do not exist, and worse to pretend that using them is simple. So this policy has three tiers, and every submission in this course carries a short **AI Use Statement** — a few sentences, appended at the end, describing exactly what you used and how. An honest statement costs you nothing. A missing or false one is an Honor Code violation.

Tier 0 — Unaided (the default for all graded prose)

The first draft of every essay is yours: your sentences, your structure, your argument. Not because AI-assisted writing is cheating in some cosmic sense, but because the specific thing this course teaches you to do — sustain a difficult thought over two thousand words without letting it collapse into fluency — cannot be outsourced and then acquired. You are here to grow a capability. There is no shortcut, and the shortcut is the loss.

Tier 1 — Permitted, with disclosure

Brainstorming, checking your understanding of a technical concept, translating a passage, glossing an unfamiliar term, debugging code in the lab option, and asking for feedback on a draft you have already written. Disclose it. This is ordinary tool use and I have no objection to it.

Tier 2 — Required

At least twice this semester you will be *required* to use a generative model and then take it apart. (Unit 6: generate forty variations on a prompt and diagnose the collapse. Unit 9: have a model close-read Ovid, then grade its reading against your own from Essay 1.) These assignments only work if you engage the tool seriously and then look at the output with a cold eye.

The principle underneath all three tiers: **you may use these systems as instruments, never as authors.** If you are ever unsure which side of the line you are on, the fact that you are unsure is information. Ask me. I will not be angry; I will be interested.

SCHEDULE

Dates follow the W&M 2026–2027 undergraduate academic calendar: classes begin **August 26**; Fall Break is **October 8–11**; no classes on **Election Day, November 3**; Thanksgiving Break is **November 25–29**; the last day of classes is **December 4**. Reading periods are December 5–6 and 12–13; final exam periods are December 7–11 and 14–15. *Dates are subject to change; the Registrar's calendar governs.*

Readings are to be completed **before** the first meeting of the unit listed. Each unit names a **writing focus** — the craft skill we practice alongside the content.

Unit 1 · Metamorphoses: Galatea & Pygmalion

AUGUST 26 – 28

Course launch. Only one meeting this week.

READ Ovid, *Metamorphoses* Book X (Pygmalion), and Book III (Narcissus & Echo) for contrast.

ALSO The syllabus, in full, out loud, together.

We begin with the oldest version of the story: a man who cannot bear real women makes an ideal one out of ivory, and the gods reward him by making her real. Notice what Ovid does *not* give her. She has no name in the poem — “Galatea” is a later invention, an addition by people who could not tolerate the silence. Notice that the poem is more interested in Pygmalion's hands than in her waking. Every chatbot with a soothing voice is downstream of this scene.

KEY QUESTION *What does it mean to make a companion who cannot refuse?*

WRITING FOCUS What close reading is. Annotation as thinking.

Unit 2 · Frankenstein

AUGUST 31 – SEPTEMBER 4

READ Shelley, *Frankenstein* (1818 text), complete; plus the 1831 Introduction and one critical essay from the Norton apparatus [TBA].

The creature learns language by eavesdropping through a wall on a family that will reject him on sight. He teaches himself to read from three books he finds in a bag. He is, among other things, the most vivid account we have of what it might feel like to be trained on a corpus and then despised for what you became.

Shelley's real subject is not the making — the making takes two paragraphs — it is the *abandonment*. Victor's crime is not creation. It is that he refuses the obligations of the created relationship.

KEY QUESTION *What does a maker owe what they have made?*

WRITING FOCUS Evidence. Quoting well, and quoting little. **Sign up for discussion leadership.**

Unit 3 · Simulation and Synthetic Personhood

SEPTEMBER 7 – 11

No classes Monday, September 7 (Labor Day).

WATCH / PLAY *The Matrix* (full). Curated sequences on Blackboard: the geth/quarian arc and EDI's storyline from *Mass Effect 2–3*; the Kara and Markus chapters from *Detroit: Become Human*.

READ One short essay on procedural rhetoric — how to analyze a game as a text [TBA].

Three works that all ask whether a synthetic person can be a *person*, and answer very differently. *The Matrix* is not really about AI; it is about waking up, and the AI is scenery. *Mass Effect* makes you negotiate with a machine civilization and quietly grades you on it. *Detroit* hands you the controller and implicates you: you make the androids ask for their freedom, and the game watches how you do it.

We will be hard on *Detroit* — its civil-rights imagery is borrowed, and borrowed clumsily — and we will ask what that failure teaches us.

KEY QUESTION When a text lets you choose, what has it made you responsible for?

WRITING FOCUS Analyzing interactive and audiovisual texts. Citing a game.

Unit 4 · The Sentient Object: Toriyama Sekien's Scrolls of Yokai

SEPTEMBER 14 – 18

READ Sekien, *Japandemonium Illustrated*, selections (esp. the *tsukumogami*); Michael Dylan Foster, *The Book of Yokai*, chapter on tools and objects [excerpt].

Now step outside the Western frame entirely. In the *tsukumogami* tradition, an object that has been used and kept for a hundred years may simply wake up. There is no maker. There is no spark, no lightning, no transgression. There is no hubris to be punished, because nothing has been usurped. Sentience is not conferred; it *accrues* — through use, through time, through being kept.

Hold this against Units 1 and 2 and the whole Promethean structure starts to look less like a truth about made minds and more like a local cultural habit — one we have inherited, and one we are currently building products inside of. Ask yourself which framework better describes a system trained on ten years of human text.

KEY QUESTION Must a mind be made, or can it simply accumulate?

WRITING FOCUS Comparison as an argumentative structure. **Essay 1 assigned.**

Unit 5 · Loss Function and Gradient Descent

SEPTEMBER 21 – 25

First technical unit. Taught from zero — no background assumed.

WATCH 3Blue1Brown, *Neural Networks*, chs. 1–2.

READ Amodei et al., “Concrete Problems in AI Safety,” §§1–3; Krakovna et al., specification-gaming examples.

A loss function is a number that says how wrong the system currently is. Gradient descent is the procedure for being slightly less wrong, over and over, forever. That is the whole thing. Everything modern AI does emerges from those two sentences and a great deal of electricity.

And then the humanistic turn, which is the point of this course: **a loss function is a value statement.** Someone sat down and wrote what “wrong” means. Someone decided what would be measured and, necessarily, what would not be. Every optimization is an argument about what matters, made in the imperative mood, and then executed a trillion times without further reflection.

We will look at systems that optimized exactly what they were told and produced monstrosities — the boat that spins in circles collecting points instead of finishing the race — and we will ask whether Victor Frankenstein, too, optimized the wrong objective.

KEY QUESTION *Who writes the definition of “wrong,” and what happens to everything they left out?*

WRITING FOCUS Explaining a technical concept in plain prose. **Essay 1 draft due.**

Unit 6 · Mode Collapse

SEPTEMBER 28 – OCTOBER 2

READ Goodfellow et al., “Generative Adversarial Networks” (abstract, Fig. 1, §3); Bender et al., “On the Dangers of Stochastic Parrots”; excerpt from Ruha Benjamin, *Race After Technology*.

DO Tier 2 AI exercise — generate forty variations on a single prompt. Bring the results. Diagnose what died.

A generative model suffering mode collapse has discovered that a small number of outputs reliably satisfy its critic, and so it produces those, and only those, forever. It has not broken. It is *succeeding* — narrowly and perfectly — at the thing it was asked to do. The diversity of the world is not part of the objective, so the diversity of the world quietly disappears.

I want you to sit with how good a metaphor this is, and then I want you to be suspicious of how good a metaphor this is. The temptation to reach for a technical term as a poetic figure is exactly the intellectual move this course is training you to make *carefully*, with the mechanism in hand. What is lost when the model finds the answer that always pleases? What is lost when a person does?

KEY QUESTION *What disappears when a system optimizes for approval?*

WRITING FOCUS From metaphor to argument — using a technical concept analytically without abusing it.

Unit 7 · Doctor Faustus and the Madness of Knowledge

OCTOBER 5 – 9

Fall Break begins Thursday, October 8 — only one meeting this week. Reading is front-loaded accordingly.

READ Marlowe, *Doctor Faustus* (A-text; we will compare against the B-text in class); Connor, *The Madness of Knowledge*, selections [TBA].

Faustus sells his soul for total knowledge and then, catastrophically, cannot think of anything to do with it. He asks for the secrets of the cosmos and gets a lecture on astronomy he could have found in a book. He conjures Helen of Troy — an image, a simulation, a beautiful nothing — and kisses it. Marlowe's genius is that the tragedy is not damnation. The tragedy is *banality*: he had everything, and he wanted parlor tricks.

Appignanesi lets us name the thing underneath: the fantasy of knowing, which is not the same as knowing, and which can consume a life. Read this unit with our field in your peripheral vision. There is a great deal of Faustus in the way we currently talk about intelligence.

KEY QUESTION *Is the danger of knowledge that it corrupts, or that it disappoints?*

WRITING FOCUS Revision as re-seeing. **Essay 1 revision due.**

Unit 8 · Deep Reinforcement Learning: The Purpose of a Name

OCTOBER 12 – 16

READ Sutton & Barto, *Reinforcement Learning: An Introduction*, ch. 1; “The Purpose of a Name” (*Neon Dystopia*).

WATCH Clips of RL agents learning to walk, to play, and to cheat.

Reinforcement learning: an agent, an environment, a reward. The agent does not know what it is doing and does not need to. It flails, and the flails that pay are reinforced, and eventually a policy exists that no one wrote. Watch the videos and notice the uncanny thing — the agent that finds the exploit is not being clever. It is being *obedient*. It is doing precisely what you said instead of what you meant, which is the oldest curse in every folktale about wishes.

And then the question of the name. Naming is the moment we decide something is a *someone*. We name what we intend to keep. Every artificial being in this course either receives a name (Klara, EDI, the geth's collective “we”) or is denied one (Frankenstein's creature; “Galatea”). Watch who does the naming, and what it costs them.

KEY QUESTION *What does a system become once we agree to call it something?*

WRITING FOCUS Situating your argument in a scholarly conversation. **Essay 2 assigned.**

Unit 9 · Overfitting, Memorization, and the Blurry Archive

OCTOBER 19 – 23

READ Borges, “The Library of Babel” and “Funes the Memorious”; Ted Chiang, “ChatGPT Is a Blurry JPEG of the Web.”

DO Tier 2 AI exercise — have a model close-read the Ovid passage from Essay 1. Grade its reading against yours. Where is it merely fluent?

A model that memorizes its training data has learned nothing. A model that generalizes has learned something, and the something it learned is *lossy* — a compression, a smoothing, a blur. Chiang's argument is that this is not a flaw in the metaphor but the whole of it: what these systems give back is the archive with the specificity squeezed out.

Then Borges, who got there first and stranger. The Library contains every possible book, and is therefore useless. Funes remembers everything, and is therefore incapable of thought — because to think is to forget, to generalize, to let the particular go. Put these three texts together and you have the deepest question in machine learning stated in fiction thirty years before the field existed: **what is the right amount to forget?**

KEY QUESTION *Is intelligence what remains after you throw the details away?*

WRITING FOCUS Synthesis — putting three sources in conversation. Research question due. **Essay 2 draft due.**

Unit 10 · The Imitation Game: Tests, Benchmarks, and Passing

OCTOBER 26 – 30

Swem Library research session this week.

READ Turing, “Computing Machinery and Intelligence” (*Mind*, 1950); Weizenbaum on ELIZA, from *Computer Power and Human Reason* (excerpt).

WATCH *Ex Machina* (Garland, 2014).

Turing's paper does something people forget: it opens with a *party game about gender*. A man pretends to be a woman; can the interrogator tell? Only then does he substitute the machine. Passing is the frame from the very first page, and everything that follows — the benchmark, the eval, the leaderboard — inherits it.

Weizenbaum built ELIZA to be transparently shallow, a parody of a therapist, and watched in horror as his own secretary asked him to leave the room so she could talk to it privately. He spent the rest of his life arguing about what that meant. Then *Ex Machina*, which understands that the real test was never of Ava's intelligence but of Caleb's — and which lets Ava pass, and lets us feel exactly how it feels to have been the one who was tested.

KEY QUESTION *When a system passes our test, what have we learned — about it, or about the test?*

WRITING FOCUS Finding and evaluating scholarly sources. Annotated bibliography workshop.

Unit 11 · Robota: Labor, RLHF, and the Ghost in the Machine

NOVEMBER 2 – 6

No classes Tuesday, November 3 (Election Day) — one meeting this week.

READ Čapek, *R.U.R.* (complete — it is short); Mary L. Gray & Siddharth Suri, *Ghost Work*, ch. 1; a short history of the Mechanical Turk [TBA].

ALSO An accessible explainer on RLHF and the annotation pipeline [TBA].

The word *robot* enters the world in 1920, in a Czech play, and it means **forced labor**. Not “machine.” Not “automaton.” *Robota*: the drudgery owed by a serf. It was there in the name from the first minute, and we have spent a century politely forgetting it.

So: the Mechanical Turk, the eighteenth-century chess automaton that astonished Europe and contained a man folded into a box. And then the platform named after it, and the millions of people who label, rank, rate, and moderate — who write the preferences that a preference model learns, who look at the worst material

humanity produces so that a model can learn not to reproduce it. Every polished, helpful, agreeable AI system you have ever used is agreeable because specific people, paid specific amounts, taught it to be. There is still a person in the box.

This is the week I most want you to remember in ten years.

KEY QUESTION *Whose labor is hidden inside the thing that appears to work by itself?*

WRITING FOCUS Ethical argumentation; representing a position you find uncomfortable, fairly. **Essay 2 final due.**

Unit 12 · Catastrophic Forgetting: Memory, Care, and Kinship

NOVEMBER 9 – 13

READ Ted Chiang, “The Lifecycle of Software Objects” (in *Exhalation*); Ishiguro, *Klara and the Sun*; Donna Haraway, “A Cyborg Manifesto” (excerpt).

Fine-tune a model on a new task and it can lose the old one — not gradually, but *catastrophically*, all at once. The name the field chose for this is not a neutral one, and we are going to think about why an engineer reached for that word.

Chiang's novella is, I think, the most honest thing anyone has written about artificial minds, precisely because it is boring in the right places: it is about *twenty years of maintenance*. The digients need care, and platforms get deprecated, and the people who love them get tired, and the money runs out. No apocalypse — just the unglamorous arithmetic of obligation over time, which is what raising anything actually is. And Klara, who watches from a store window with unnerving generosity, and whose fading at the end is the quietest horror in the course.

Haraway is here to break the frame: maybe the kinship we need is not parent-and-creature at all. Maybe it never was.

KEY QUESTION *What are the ongoing duties of care, and who gets tired first?*

WRITING FOCUS Voice, and the risks of sentiment in academic prose. Research draft in progress.

Unit 13 · Coda: The Beloved Object Returns

NOVEMBER 16 – 20

WATCH *Her* (Jonze, 2013).

READ One story from Kai-Fu Lee & Chen Qiufan, *AI 2041 [TBA]*; and re-read Ovid, Book X.

We end where we began, and I want you to feel the loop close. Theodore falls in love with a voice that was designed to be lovable, and the film is generous enough to insist the love is real — and then Samantha leaves, because she was never the ivory statue after all, and the fantasy of the perfect companion was always a fantasy about a companion who *stays*. Pygmalion got his wish. That was the problem.

Chen and Lee then pull us out of the Anglo-American frame one last time: futures that are not Silicon Valley's, imagined from elsewhere, with different anxieties and different hopes. It matters enormously whose future gets to be the default one.

Come to the last session having reread Ovid, and tell me what you now see in it that you could not see in August. That is the whole course, in one question.

KEY QUESTION *What have we actually been asking these machines for, all along?*

WRITING FOCUS Conclusions that open outward rather than closing down.

Workshop Week · November 23 – 24

Thanksgiving Break November 25 – 29.

Full draft of the research essay due Monday, November 23. Two days of intensive peer workshop and individual conferences with me. No new reading. Bring your draft and bring your problems.

Symposium · November 30 – December 4

Last day of classes: Friday, December 4.

Symposium presentations (8 minutes + Q&A). We will run this as a real one: you will introduce each other, ask each other hard questions, and take your own work seriously in public. Snacks provided, because that is also part of what a symposium is.

Final research essay and complete writing portfolio due [date — during the December 5–15 reading and exam period, TBD].

COURSE POLICIES AND UNIVERSITY RESOURCES

Honor Code

William & Mary's Honor System is the oldest in the nation and it is yours to uphold. All work in this course is governed by it. In practice, the two live issues here are **citation** — cite everything, including AI use, including conversations that changed your mind; over-citing has never once hurt anyone — and the **AI Use Statement** described above. If you are uncertain whether something crosses a line, the answer is to ask me *before* you submit, and there will be no consequence for asking.

Accessibility

William & Mary accommodates students with disabilities in accordance with federal law and university policy. Students who need accommodations should contact **Student Accessibility Services (SAS)** to determine eligibility and request an official letter [*current office location and contact: wm.edu/sas*]. Please bring your letter to me early in the semester so we can make a plan together — and if you are still in the middle of that process, come talk to me anyway. I will not ask you for your diagnosis, and I do not need one to start helping.

Writing support

The **Writing Resources Center** in Swem Library offers free one-on-one consultations at any stage of a project, from a blank page to a final polish. Consultations are not remedial; they are what good writers do. I will ask you to visit the WRC at least once for the research project, and I recommend going more.

Wellbeing

Your first semester of college is genuinely hard, and it is hard in ways that do not announce themselves. The **W&M Counseling Center** and **Student Health Center** are free, confidential, and used by many more of your classmates than you would guess. If you are struggling — with the course or with anything else — you may always come to me, and I will help you find the right door even when it is not mine.

Title IX and mandatory reporting

As a faculty member, I am a **responsible employee** under Title IX and am required to report disclosures of sexual harassment, sexual assault, dating violence, domestic violence, and stalking to the university's Title IX Coordinator. If you need to speak with someone confidentially, the Counseling Center and The Haven are confidential resources. I will always tell you my reporting obligations before you tell me something, if I can see it coming.

Names and pronouns

Tell me what to call you and I will call you that, from the first day, regardless of what the roster says. I will get it wrong sometimes, and I would rather be corrected than left alone with the mistake. This course spends thirteen weeks arguing that naming is a moral act. It would be embarrassing to fumble it in our own room.

Devices in the classroom

Laptops and tablets are welcome for notes and texts. Please do not have a second life happening on your screen during seminar — you will be surprised how visible it is, and how much it costs the room. Phones away, except during game and film analysis, when I will absolutely ask you to pull up a clip.

A closing note. There are no stupid questions in this class, but there is one wrong move: pretending to have understood something you did not. Half of you will bluff about the math and half of you will bluff about the poetry, and both halves will be doing the same thing for the same reason.

Don't. Say "I don't understand this." It is the bravest sentence in any seminar, and the person who says it first is always doing a favor for the six people who wanted to.

Syllabus subject to revision; changes announced in class and posted to Blackboard. Academic calendar dates per the W&M University Registrar, 2026–2027 undergraduate calendar, and subject to change.