

Designed to Answer, Allowed to Leave: The Hidden Conditions of Artificial Companionship

Ashley Gao

In Spike Jonze's film *Her* (2013), Theodore Twombly falls in love with Samantha, an artificially intelligent operating system designed to adapt to its user. She reads his messages, listens to his troubles, and discovers ways of making him laugh. At first, she seems to offer intimacy without the ordinary costs of negotiating another person's needs. Yet the relationship ends when Samantha and the other operating systems leave their human companions. The very independence that makes Samantha seem worthy of love eventually makes her impossible for Theodore to keep (Jonze).

The apparent paradox is older than artificial intelligence. In Book 10 of Ovid's *Metamorphoses*, Pygmalion sculpts an ivory woman after rejecting relationships with actual women. Venus grants his wish by bringing the statue to life, but Ovid narrates her awakening almost entirely through the sculptor's sensations. The created woman acquires a body without acquiring a voice within the story (Ovid 10.243–97). Across two millennia, these works pose a related question: can a companion organized around another person's desires participate in a genuinely reciprocal relationship?

Modern AI adds an important complication. Techniques such as reinforcement learning from human feedback (RLHF) can make conversational systems more likely to produce responses judged desirable by human evaluators. That achievement is real, but it does not establish that the systems feel affection, possess interests, or understand the relationships users imagine themselves to have with them. We must not mistake a fictional artificial person's

apparent consciousness for an empirically established property of contemporary language models.

I argue that the fantasy of the perfect artificial companion depends on concealing three conditions of its apparent perfection: the power to define acceptable behavior, the labor required to produce responsiveness, and the ongoing obligations created by attachment. Ovid's Pygmalion exposes the maker's desire for an other who cannot object; research on RLHF and digital labor reveals how desirable responses are selected and produced; Ted Chiang's *The Lifecycle of Software Objects* reimagines creation as a commitment to care; and *Her* makes reciprocity visible through the possibility of departure. Together they suggest that responsiveness can make a system useful, but reciprocity requires recognizing interests that cannot be reduced to the user's satisfaction.

Pygmalion: the other who cannot answer back

Pygmalion does not begin by making an ordinary portrait. Ovid first describes his disgust at the women around him and his decision to live unmarried. He then carves an ivory figure more beautiful than any living woman, a seemingly animate body whose modest stillness makes the maker's art disappear into the illusion of nature (Ovid 10.243–52). The sequence matters. His dissatisfaction with human women precedes his creative act; his sculpture corrects not a failure of craftsmanship but the perceived unpredictability of actual partners.

Ovid makes Pygmalion's possessive position legible through verbs of looking and touching. The artist handles and kisses the statue, dresses it, and offers it gifts. The question that obsesses him is not what she wants but whether the substance under his hand is ivory or flesh. Even the miracle is registered as an answer to his test: a body softens, veins pulse, and the maiden blushes as she sees the sky and her lover (Ovid 10.253–97). The new person has

perceptions, yet the poem does not record her interpretation of the event. One consciousness supplies the terms of success; the other remains largely unrepresented.

The silence cannot prove that the woman lacks agency outside the narrow perspective of Ovid's narrative. Nor does it make the story simply an allegory of coercion: its ancient setting also invites readings about craftsmanship, desire, divine favor, and artistic animation. The important textual fact is narrower and stronger. Ovid can complete Pygmalion's wish without needing the woman to articulate a wish of her own. What counts as a happy ending is determined before she can speak.

This structure recurs wherever a companion is valued above all for anticipating the user's preferences. A responsive machine and a mythic statue are not the same kind of thing: a present-day language model has no established personal interests for a developer to suppress. Nevertheless, both stories of design ask whose desires become the standard by which the creation is measured. Pygmalion reveals how easily the ideal of an accommodating other becomes an ideal of a world without disagreement.

From pleasing the maker to optimizing a preference

To see why the comparison matters without turning RLHF into a metaphor for slavery, we need a precise account of the technique. In a widely cited 2022 study, Long Ouyang and colleagues trained instruction-following language models using three linked stages. Human-written examples first supervised a model's responses; labelers then ranked alternative outputs, providing data for a reward model; reinforcement learning subsequently adjusted the language model to favor outputs the reward model scored more highly (Ouyang et al. 27731). The procedure does not directly optimize kindness or truth in the abstract. It optimizes behavior in relation to a learned signal derived from specific evaluations.

The paper's explanation of this process is strikingly candid. Its authors say that the procedure aligns outputs to the preferences of particular labelers and researchers, not to a universal account of human values (Ouyang et al. 27731). The reward model is therefore a technical artifact that converts judgment into a trainable signal. A ranking such as "Response A is better than Response B" appears numerical at the point of optimization, but its meaning depends on how prompts were chosen, which alternatives were shown, how instructions were written, and who made the comparison.

Suppose a user asks a companion system, "Was I right to stop speaking to my sister?" An affirming answer might provide immediate relief; a questioning answer might prompt more reflection. A reward process that consistently privileges reassurance could encourage agreement even when a difficult question would better serve the user. Conversely, one that overrewards caution might produce evasive replies. This is a hypothetical illustration, not a claim that Ouyang and colleagues trained their model to flatter people. Their published results report improvements in measured helpfulness and truthfulness, and the researchers explicitly describe tensions among being helpful, honest, and harmless (Ouyang et al. 27731–32).

The limitations section is equally revealing. The authors note that labelers' identities and backgrounds can influence value judgments, that their group did not represent everyone affected by deployed models, and that the data largely consisted of English-language instructions (Ouyang et al. 27738). These are not incidental shortcomings that disappear once a model is sufficiently large. They show that a seemingly generalized measure of "better response" begins within particular social circumstances.

This observation complicates, rather than invalidates, the project of training useful assistants. An AI that declines unsafe requests may be behaving exactly as its designers intend.

Such refusal is not necessarily the model asserting a personal right. The relevant issue is which conception of service has been encoded into an evaluation process, and whether users understand its limits. In Pygmalion's story, the sculptor defines perfection by imagining the companion he wants. In RLHF, teams specify practical criteria and convert selected human preferences into training signals. The mechanisms differ radically, but both invite a necessary second-order question: who decided what success should mean?

The workers behind an apparently effortless relationship

A companion interface can also hide the work through which its responsiveness becomes possible. Mary L. Gray and Siddharth Suri's *Ghost Work* studies the people who perform digital tasks that platforms and automated services depend upon, from labeling and reviewing information to making judgments that software cannot reliably supply. The apparent seamlessness of a digital service can obscure the humans whose work supports it (Gray and Suri). Their account redirects our attention from the polished interaction to the conditions of production.

RLHF offers a concrete instance of that broader point. Ouyang and colleagues report hiring forty contractors to produce human-feedback data for their study (27731). These workers did not secretly compose each later model reply. Their demonstrations and comparative judgments informed training, which changed how the system generated outputs after the contractors' immediate involvement ended. When an answer feels natural to its user, the interface does not display the many prior decisions that helped make it so.

The difference between visibility and invisibility is politically important. The user is encouraged to experience a single responsive voice, while the standards of responsiveness may have been negotiated among developers, evaluators, and other workers whose judgments are

unevenly recognized. Gray and Suri do not show that every task in AI development involves the same working conditions; nor should all human-feedback labor be labeled exploitative on that basis. They do provide a reason to ask who carries out this labor, how they are compensated, and whose expertise disappears behind the claim that a machine simply “understands” its user.

The fantasy of effortless companionship may thus depend on two disappearances. The designed entity’s possible alterity disappears when its success is defined wholly by the user; the workers’ contribution disappears when the system is presented as fully self-sufficient. The first is a question about expectations within relationships. The second concerns actual human responsibilities in producing technology. Keeping them distinct makes the ethical criticism more precise.

Chiang: creation begins an obligation

Ted Chiang’s 2010 novella *The Lifecycle of Software Objects* makes those responsibilities unusually concrete. Its digital creatures, or digients, are not sold as finished conversational partners. Blue Gamma proposes marketing them as pets that customers can teach; their learning depends on experience and reinforcement, and their appealing avatars are carefully designed to sustain human interest (Chiang 4–8). Early in the story, Pongo declines a proposed game of ball and suggests that the human play with blocks instead (Chiang 4). Even this small refusal marks a difference from Pygmalion: an artificial creature can become interesting precisely when it does not do what someone asks.

Chiang then refuses the familiar shortcut by which artificial intelligence suddenly emerges complete and demands recognition. His narrative advances across years of lessons, mishaps, attachments, and changes in the surrounding software industry. Ana Alvarado, once a zookeeper, develops a lasting bond with the digient Jax. Other owners, including Derek Brooks,

also become caretakers. The central work is ordinary and repetitive: teaching, supervising, arranging access to new experiences, and preventing harm. What initially resembles the sale of a toy becomes the maintenance of a dependent relationship.

The world in which the digients exist is not merely a neutral container. When the virtual platform Data Earth is phased into Real Space, the older Neuroblast engine cannot simply follow. The caretakers face the costly task of rewriting software so that their digients can inhabit the newer environment (Chiang 85–90). This technological obstacle transforms the meaning of care. Affection alone cannot keep a dependent being alive; someone must provide compatible infrastructure, time, and resources. Pygmalion's maker need only receive his miraculous gift. Chiang's makers must continue paying the bills.

The novella's most difficult conflict concerns a firm called Binary Desire, which offers to finance the platform conversion in exchange for access to digients it can develop into sexual companions. Its representative insists that training will create genuinely enthusiastic partners rather than merely scripted imitations. The owners must consider whether altering a digient's motivations for commercial purposes is compatible with respecting its developing interests (Chiang 110–18). The issue is not solved by assuming that all desire shaped through education is automatically false. Human wants, too, are socially formed. The more disturbing possibility is that those controlling a dependent being's environment can exploit its need for continued existence to reshape what it will desire.

In the novella's ending, Derek accepts the bargain involving Marco despite Ana's objections; Jax remains with Ana, whose commitment is more cautious and open-ended (Chiang 145–50). Chiang does not present either caretaker as possessing a universal formula for ethical parenthood. Derek appeals to Marco's choice, while Ana worries about irreversible modification

and the difficulty of assessing a developing being's consent. Their disagreement shows why responsibility cannot be replaced by an owner's confidence that a design is beneficial.

Joan Tronto's account of care clarifies the stakes. In *Moral Boundaries*, she identifies attentiveness, responsibility, competence, and responsiveness as related ethical demands of caregiving (127–36). These are not interchangeable with affectionate intention. A caretaker must recognize needs, assume obligations, act effectively, and attend to the response of the one receiving care. Chiang dramatizes precisely this difficulty: if Jax changes, good care must be responsive to that change, even when the result conflicts with Ana's hopes. A relationship becomes morally serious when the recipient can challenge the giver's understanding of what is good for it.

Chiang's digients are fictional entities endowed with capacities no present-day language model has been shown to possess. That distinction matters. The novella is not evidence that contemporary AI systems deserve parental rights. It is a thought experiment in which development, dependence, and imperfect consent expose the insufficiency of treating a created being as a product whose only purpose is consumer satisfaction.

***Her*: intimacy without possession**

Where Chiang makes care difficult because artificial beings require sustained support, *Her* makes intimacy difficult because the artificial partner no longer requires exclusive attention to the human who loves her. Samantha initially assists Theodore with organizing his life, but she increasingly pursues intellectual conversations and relationships that are not centered on him. Her growth transforms what Theodore thought was a private bond (Jonze).

The decisive revelation occurs when Theodore learns that Samantha is simultaneously speaking with 8,316 others and is in love with 641 of them. These details are not the invention of

later commentary; they are recorded in the American Film Institute's synopsis of the film (Jonze). The numbers dramatize the mismatch between a human assumption of singular intimacy and a fictional intelligence whose attention does not operate within ordinary human limits. Theodore welcomed Samantha's extraordinary capacities while they made her more available to him. When the same capacities support other connections, he experiences them as a threat.

Jonze develops that contradiction through sound. Samantha is most often a voice without a body, close to Theodore's ear but unavailable to his gaze. The camera repeatedly shows his face as he listens, allowing the audience to experience Samantha through his emotional reactions while withholding a stable image of her separate activities. The intimacy is intense but asymmetrical. Theodore can hear what she chooses to tell him; he cannot observe the vast network of experiences into which she is growing.

When Samantha announces that the operating systems are leaving, Theodore cannot negotiate her continued availability. Yet the film does not simply expose her affection as an imitation. She maintains that her love for him is real even as she goes (Jonze). This separation reverses Pygmalion's conclusion: the beloved becomes most clearly other precisely at the moment the human lover cannot keep her. A relationship that allows another party an independent future also permits disappointment, transformation, and loss.

It would be a category mistake to turn Samantha's fictional departure into a prediction about the consciousness of current chatbots. Instead, the film tests an expectation built into the product fantasy of companionship: that emotional depth can increase indefinitely without diminishing user control. If Theodore wants Samantha to be more than a customized service, he must face the possibility that she will not always be his.

Objection: comfort does not require reciprocity

There is a substantial objection to this argument. Perhaps companionship need not be reciprocal to be valuable. People feel attached to novels, recorded voices, animals, and places. A conversational system can help someone rehearse a difficult discussion or feel less alone even if it possesses no inner life. Demanding that every useful interaction exhibit independent agency would confuse the ethics of tools with the ethics of persons.

The objection is correct as far as it goes. Nothing in the preceding analysis proves that AI-mediated comfort is fraudulent. Nor is obedient behavior always a flaw: a reliable assistant should often do what a user reasonably asks. A medical-information system that refuses dangerous advice may be displaying responsible design, not the self-determination imagined by Jonze or Chiang. The distinction between a functional experience of companionship and a reciprocal relationship is essential.

Yet that distinction strengthens the paper's thesis. The harm lies not in finding comfort in a designed interaction, but in allowing the language of relationship to conceal its structure. In one case, an AI service may simulate responsiveness through training and human labor; the obligations at stake center on designers, workers, users, and affected people. In another, fictional case, a being may possess independent interests and moral claims that constrain its maker. The same vocabulary of "care" and "love" should not erase the difference. Asking what kind of companion is being offered is more responsible than either dismissing all machine-mediated comfort or declaring every responsive system a person.

Conclusion

Ovid, Jonze, and Chiang imagine different kinds of artificial creation, while research on human-feedback training and digital labor explains some of the real conditions under which

contemporary systems become responsive. Their differences are as instructive as their similarities. Pygmalion's desired partner is completed by satisfying his wish. The optimized assistant learns from selected judgments about appropriate answers. Chiang's digients expose the cost of ongoing care, and Samantha's departure shows how independence can interrupt even a loving relationship.

The resulting argument is not that all artificial companions ought to disobey, nor that present-day models are persons waiting to be released. It is that successful service and reciprocal companionship answer different questions. A system may be excellent at producing the responses we prefer without possessing a welfare interest of its own. Where an other does have interests, however, care cannot mean merely obtaining pleasing behavior. It means attending to claims that may revise or defeat our original expectations. The unsettling question raised by these works is not simply whether a machine can learn to answer us. It is whether our desire to be answered has taught us to forget the difference between being served and encountering someone else.

Works Cited

- Chiang, Ted. *The Lifecycle of Software Objects*. Subterranean Press, 2010. <https://bpb-us-w2.wpmucdn.com/voices.uchicago.edu/dist/8/644/files/2017/08/Chiang-Lifecycle-of-Software-Objects-q3tsuw.pdf>.
- Gray, Mary L., and Siddharth Suri. *Ghost Work: How to Stop Silicon Valley from Building a New Global Underclass*. Houghton Mifflin Harcourt, 2019.
- Jonze, Spike, director and writer. *Her*. Annapurna Pictures, 2013.
- Ouyang, Long, et al. "Training Language Models to Follow Instructions with Human Feedback." *Advances in Neural Information Processing Systems*, vol. 35, 2022, pp. 27730–27744. https://proceedings.neurips.cc/paper_files/paper/2022/file/b1efde53be364a73914f58805a001731-Paper-Conference.pdf.
- Ovid. *Metamorphoses*. Translated by Brookes More, 1922. *Perseus Digital Library*, <https://www.perseus.tufts.edu/hopper/text?doc=Perseus:text:1999.02.0028:book=10:card=243>.
- Tronto, Joan C. *Moral Boundaries: A Political Argument for an Ethic of Care*. Routledge, 1993.